

SPECRUN: The Danger of Speculative Runahead Execution in Processors

Chaoqun Shen
Hunan University, China
shencq@hnu.edu.cn

Gang Qu
University of Maryland, USA
gangqu@umd.edu

Jiliang Zhang*
Hunan University, China
zhangjiliang@hnu.edu.cn

Abstract

Runahead execution is a continuously evolving microarchitectural technique for processor performance. This paper introduces the first transient execution attack on the runahead execution, called SPECRUN, which exploits the unresolved branch prediction during runahead execution. We show that SPECRUN eliminates the limitation on the number of transient instructions posed by the reorder buffer size, enhancing the exploitability and harmfulness of the attack. We concretely demonstrate a proof-of-concept attack that causes leaking secrets from a victim process, validate the merit of SPECRUN, and design a secure runahead execution scheme. This paper highlights the need to consider the security of potential optimization techniques before implementing them in a processor.

Keywords: Microarchitecture Security, Runahead Execution, Transient Execution Attack, Spectre

ACM Reference Format:

Chaoqun Shen, Gang Qu, and Jiliang Zhang. 2024. SPECRUN: The Danger of Speculative Runahead Execution in Processors. In *61st ACM/IEEE Design Automation Conference (DAC '24)*, June 23–27, 2024, San Francisco, CA, USA. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/3649329.3655932>

1 Introduction

Microarchitectural techniques are the crucial foundations for improving modern processors' performance and power efficiency. Despite significant advancements in those metrics, the design of ISA is treated with little security considerations. As demonstrated by Spectre [1] and Meltdown [2], the CPU microarchitecture is susceptible to various vulnerabilities. Attackers may exploit subtle microarchitectural states to extract privileged data [3], obtain indirect information [4], and even access secret keys protected by hardware-enforced primitives such as Intel SGX [5]. Even more alarming is that the slowing of Moore's Law has led to processors with accelerated utilization of more optimization mechanisms to improve performance continually [6]. However, as evidenced in [6–8], some promising techniques also plunge processors into a security crisis.

This work is the first to focus on analyzing a special optimization mechanism, runahead execution [9]. It is designed to tolerate long-latency loads that have to be serviced by memory. When memory stalls at the head of the reorder buffer (ROB), the execution enters

“runahead mode”, aiming to speculatively issue multiple independent memory accesses to improve memory-level parallelism (MLP). Since the initial proposal [9], several improvements [10–13] have been emerged. Runahead execution has been granted a patent by Intel and has been given the Test of Time Award. Therefore, as one of the most promising techniques for addressing the memory wall problem [10], runahead execution is expected to be deployed in future processors. However, prior research on the runahead execution has primarily focused on the advancement of performance, with little attention paid to its security.

In this work, we introduce a new attack called SPECRUN, which exploits unresolved branch prediction during runahead execution. We first implement a microarchitectural prototype of runahead execution in a processor simulator. Then, we utilize the designed proof-of-concept (PoC) attack to launch SPECRUN from unprivileged application code, which successfully leaks secret data from a victim's process. Moreover, Spectre and its variants [1, 14] have an inherent limitation in that the number of transient instructions an attacker can execute depends on the size of ROB. Thus, the attacker cannot execute overflowing gadget instructions due to the limited space. However, SPECRUN eliminates the limitation of ROB size and allows the attacker to execute more gadgets, thereby enhancing the hazard of existing attacks. At last, we propose a secure runahead execution scheme as a potential solution.

Our contributions are summarized as follows:

- **Security vulnerability.** We analyze the security of runahead execution for the first time and propose SPECRUN, a novel attack that exploits the interaction between runahead execution and branch prediction. When mixed microarchitectural mechanisms co-working on the same platform, processors may enter an inconsistent state, posing an additional challenge to system security.
- **High applicability and unique advantage.** The mixed optimization mechanism is compatible with different runahead execution and branch prediction mechanisms. Additionally, SPECRUN eliminates the limitation posed by the size of ROB on the severity of Spectre-type attacks, enabling more gadgets to be exploited in practical scenarios.
- **Secure runahead execution Mechanism.** Considering the working principle of runahead execution and the factors contributing to data leakage, we design a secure runahead execution scheme.

2 Background

2.1 Runahead Execution

When a long-latency load (called stalling load) causes the instruction window to fill up and halt the pipeline, the processor triggers runahead execution. In this process, the stalling load and its dependent instructions are marked invalid using an INV bit. These marked instructions are then removed to unblock the pipeline, which

*Corresponding author: Jiliang Zhang.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

DAC '24, June 23–27, 2024, San Francisco, CA, USA

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-0601-1/24/06

<https://doi.org/10.1145/3649329.3655932>

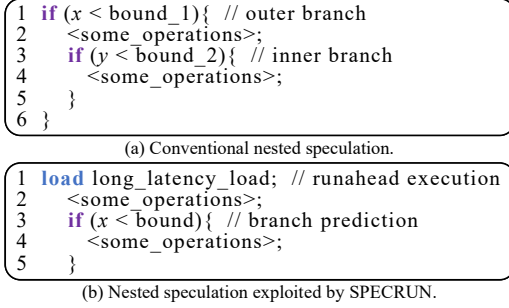


Figure 1. The sample code for nested speculative execution.

allows the processor to execute subsequent independent loads speculatively and expose the MLP. Upon the stalling load returns, the processor exits runahead mode and recovers the relevant architectural state. Due to the accurate prefetch requests generated by runahead execution, significant performance benefits are achieved.

The execution of branch predictors warrants particular attention. Take the original runahead execution [9] as an example. The way branches are predicted and resolved in runahead mode is the same as in normal mode, except for one difference: branches with INV source will never complete resolution.

2.2 Nested Speculative Execution

Branch instructions constitute a significant portion of the program due to their frequent usage. As a result, the speculative execution path frequently encounters multiple branch instructions. The size of the speculative execution window (SEW) can range from tens to hundreds of clock cycles, depending on the triggering event. Modern processors employ nested speculation to avoid front-end pipeline stalls, continuing to predict subsequent branches along the speculative path within the speculative window of the outer branch [15]. In the event that the outer branch is resolved, and a misprediction is detected, the processor discards all instructions executed along the speculative path, including the inner branches. As shown in Fig. 1, previous nested speculation referred to nested branches, while SPECTRUN utilizes a new form of nested speculation by combining runahead execution and branch prediction.

3 Threat Model

We consider a victim process and an attacker process, where the victim contains secret information that the attacker does not have direct access to but attempts to disclose. Our proposed SPECTRUN is a Spectre-type attack, and we adopt the following common assumptions for such attacks:

- The attacker is co-located with the victim on the same physical core, but she can only control one process on the target system with normal user privileges.
- The attacker cannot directly modify or access the victim's memory, but she has information on the gadget address and the target address to divulge.
- The attacker can train the target branch through her process and thus mislead the victim to perform vulnerable operations.

Since SPECTRUN is based on runahead execution, we require the victim's machine to enable the runahead technology. For research and demonstration purposes, on machines without runahead, Multi2Sim simulation can be used.

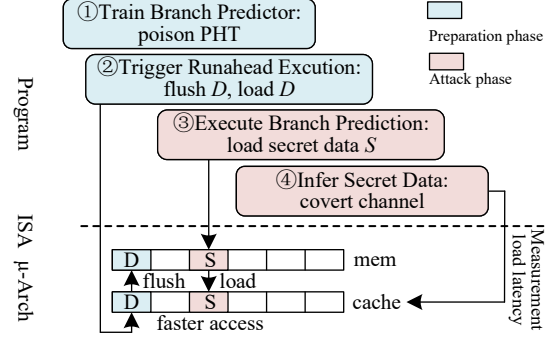


Figure 2. Overview of the workflow of SPECTRUN.

```

// Preconditions:  $D$  not present in cache
if ( $x < f(D)$ ){
    secret_data = array1[x]; //  $f(D)$  will trigger runahead
    z = array2[secret_data * N]; // access secret
}

```

Figure 3. The Code snippet for a possible gadget.

4 SPECTRUN Attack

4.1 Overview

Runahead execution improves processor performance by prefetching independent long-latency loads. Nevertheless, the speculative nature of runahead execution also introduces potential security issues. We propose a new Spectre-type attack called SPECTRUN, which exploits incorrect branch prediction within runahead execution. The attacker aims to force the victim to enter runahead mode and access sensitive data during branch prediction, ultimately leaking data through a cache covert channel.

4.2 Attack Steps

Using SpectrePHT [1] as an example, Fig. 2 outlines how SPECTRUN leaks data through the following four steps.

Step 1: Train branch predictor. The attacker runs a vulnerable branch of the victim's code cyclically. The result is that the next time the trained branch is executed, the victim will incorrectly predict the bounds check as true and execute code that should not have been executed, temporarily violating the program semantics.

Step 2: Trigger runahead execution. The operation selected to trigger runahead execution is a long-latency load, which requires the target line to have been evicted from the cache, due to a cache set contention or by employing a co-resident process to execute the *clflush* instruction. Specifically, the latter approach is adopted in this paper, wherein the attacker flushes D and then loads D .

Step 3: Execute branch prediction. The victim performs incorrect speculative execution along the training path and loads a secret-dependent memory block into the cache. Notably, the victim's branch predicate depends on the data D manipulated by the attacker. Therefore, the branch has an INV source and cannot completely be resolved during runahead execution. Fig. 3 shows an example of a potential gadget.

Step 4: Infer secret data. The victim leaves traces of secret dependencies in the microarchitecture state. Then the attacker analyses traces to recover the secret data through a covert channel.

4.3 Applicability on Runahead Execution Variants

Precise runahead [13] and vector runahead [10] are the latest variants with excellent effects. Branch prediction details are as follows:

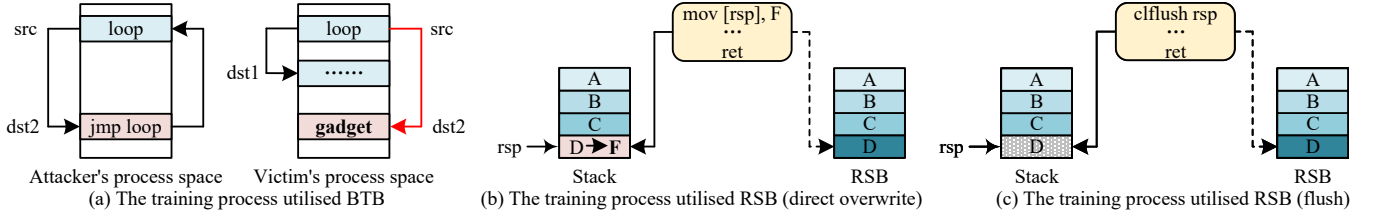


Figure 4. The mixed optimization mechanism can be applied to Spectre-type attacks. *Gadget* and *F* are associated with stalling load.

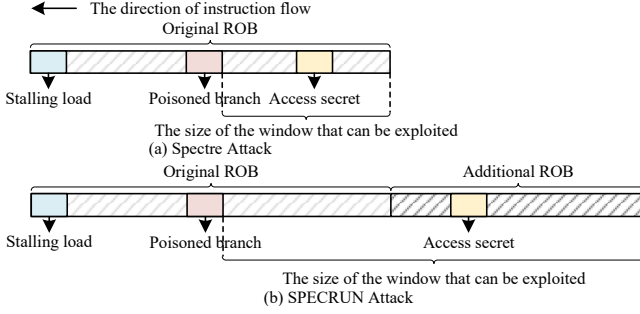


Figure 5. The difference between Spectre and SPECRUN.

Precise Runahead. Only the chains of load instructions (called stall slices) are speculatively executed. Branch instructions are not included in the stall slices because they are not involved in the computation of the load address. However, the front-end relies on the branch predictor to steer the flow of execution in runahead mode, and the branch instructions that reside in the ROB can be resolved and updated as usual. Therefore, if a branch instruction in the ROB has an INV source and is mispredicted, the corresponding stall slice guided by it can be exploited by an attacker to leak data.

Vector Runahead. To ensure the consistency of control flow, scalar branches are converted into a predicate mask for the multiple vector lanes, but only the result of the first lane is used to determine the direction of the branch. Moreover, the handling of branch instructions is similar to that in the precise runahead [9]. Hence, vector runahead is also susceptible to data leakage risks.

4.4 Applicability on Spectre Attack Variants

Once the processor enters the runahead mode, any subsequent type of branch prediction can benefit from the nested speculative execution described above. Here are the specific execution details of other Spectre attack variants [1, 14]:

SpectreBTB. As depicted in Fig.4(a), the attacker trains the virtual address congruent with *src* in his own space, which is to design a *loop* at *src* to simulate the victim's *loop*, thereby causing the branch originally directed to *dst1* to instead jump to *dst2*. Similar to SpectrePHT, we associate the gadget address *dst2* with stalling load to achieve nested speculative execution.

SpectreRSB. In the case of a direct overwrite stack, the attacker can enforce incorrect predictions by replacing the return address stored in the stack. Similarly, nested speculative execution can be enabled by associating the polluted value *F* with a stalling load. In other cases, the attacker can use *cflush* to evict the victim's stack address for simplicity, directly achieving nested speculative execution. The details are shown in Fig.4(b)-(c).

4.5 Severity of the Runahead Execution Threat

Since all changes in the program state during speculation execution must be temporarily stored in internal hardware buffers to enable

Table 1. The basic configuration of the processor.

Component	Parameter
Core	2GHz, out-of-order
Processor width	4-wide fetch/decode/dispatch/commit
Pipeline depth	6 front-end stages
Branch predictor	two-level adaptive predictor
Functional units	4 int add (1 cycle), 2 int mult (2 cycle), 1 int div (5 cycle), 2 fp add (5 cycle), 1 fp mult (10 cycle), 1 fp div (15 cycle)
Register file	80 int (64 bit), 40 fp (64 bit), 40 xmm (128 bit)
ROB	256 entries
Queue	i (40), load (40), store (40)
L1 I-cache	16KB, 4 way, 2 cycle
L1 D-cache	16KB, 4 way, 2 cycle
L2 cache	128KB, 8 way, 8 cycle
L3 cache	4MB, 8 way, 32 cycle
Memory	request-based contention model, 200 cycle

recovery in case of wrong predictions. Therefore, speculation comes to a halt once any of these buffers become full. In modern Intel CPUs, several buffers can be exhausted [16]: ROB, Branch Order Buffer (BOB), Load Buffer (LB), Store Buffer (SB), and Reservation Stations (RS), among others. However, considering factors such as buffer availability (whether the buffer has been recorded and whether the allocated entries will not be reclaimed during speculation) and universality (whether the buffer exists only in specific architectures), the ROB becomes the hardware buffer that limits the maximum number of executable instructions during speculative execution. Specifically, in Intel x86, any speculative path can contain at most as many μ ops as the entries in the ROB.

The essence of runahead execution is to convert the instruction window blocked by long-latency operations into non-blocking windows. This entails removing the stalling load from the instruction window so that subsequent instructions to be pseudo-retired from the ROB and pipeline. In other words, the runahead execution logically increases the size of ROB. Therefore, while the number of transient instructions that can be executed depends on the size of ROB, SPECRUN eliminates this limitation by exploiting the characteristic of runahead execution. SPECRUN presents a significant advantage over previous attacks because its gadgets can accommodate more non-secret-related instructions in addition to the basic instructions related to secret data load and transmission. Overall, as illustrated in Fig.5, SPECRUN reduces the requirements for gadgets and possesses a larger SEW.

5 Experimental Results

5.1 Simulation Setup

We use a runahead execution processor implemented based on the Multi2sim [17] as the attack object. The detailed configuration is listed in Table 1. The corresponding pipeline structure is shown in Fig.6. The additional hardware units highlighted in blue are intended to enable the processor to enter, execute and exit the




Figure 6. The vulnerable microarchitecture of runahead execution.




Figure 7. Standardized performance (IPC) comparison. runahead mode at the appropriate times. The components highlighted in red represent the key points of vulnerability exploited by SPECTRUN. For more details, please refer to [9]. To evaluate the effectiveness of our approach, we employed SPEC2006 as the benchmark for testing purposes. Fig.7 shows the IPC for machines with and without runahead execution. The results show that we implemented runahead execution and brought an average performance improvement of 11%. Additionally, due to the absence of specific functions in Multi2Sim required to support attack implementations, we need to add two new functions: loading data into the cache and adding a cache flush instruction.

5.2 Attack Results

Fig.8 presents a PoC example of the SPECTRUN, which accesses unauthorized memory by leveraging branch mispredictions in the victim domain during runahead mode. The victim function receives an integer x from an untrusted source. When x is less than $array1_size$, the victim accesses $array1$ with a malicious offset and uses the obtained value as an index for $array2$ (lines 3-7). To read secret data, we calculate the target address – the base address of $array1$ and then call the attacker function with a malicious index (idx). The attacker function first poisons a branch in the victim’s domain (lines 8-12). It then triggers runahead execution, causing the victim to enter runahead mode and run in nested speculative execution (lines 13-16). Finally, the attacker infers the secret data ($array1[idx]$) through the cache covert channel (lines 17-22). Note that the $array1_size$ in line 3 is associated with the stalling load to allow the victim to enter runahead mode under the attacker’s control.

Specifically, the secret data is the array index, and the attacker can infer the data loaded by the victim during execution by measuring the access time of the probe array. If the measured latency is lower than a certain threshold, we can find out that the corresponding index is speculatively loaded. Fig.9 shows the time the attacker takes to access each value in $array2$ in turn. It exhibits a significant drop at index 86, indicating that the secret data accessed by the victim is 86. The experiment emphasizes the danger of data leakage due to branch misprediction during runahead execution in the victim domain.

5.3 The Advantage of SPECTRUN

By implementing runahead execution, the processor can increase the number of ROB entries in a logical manner, thereby breaking

```

1 void victim_function (size_t x){
2   /* SpectrePHT (or other Spectre variants) */
3   if (x < array1_size) { // array1_size = f(D)
4     // mispredicted branch
5     S_value = array1[x]; // access secret
6     *entry_ptr = array2[S_value * N]; // transmit secret
7   }
8   void attacker_function (size_t idx){
9     /* training a victim branch */
10    for (int i = T; i >= 0; i--) { // T < array1_size
11      victim_function(idx);
12    }
13    /* trigger runahead execution */
14    clflush D; // set the initial state of the cache
15    load D; // enter runahead mode
16    <some_operations> // waiting for the victim's execution
17    /* probing the latency */
18    for (int j = 0; j < 256; j++) {
19      start = rdtscp();
20      array2[array1[j] * N];
21      end = rdtscp(); // measure the access time
22    }

```

Figure 8. The code snippet for SPECTRUN.




Figure 9. The probe array access time after executing SPECTRUN.

through the limitations imposed by the instruction window. Moreover, while the speculative depth of runahead execution is bounded by the available backend resources such as the Issue Queue and Physical Register File [10], research detailed in [13] has demonstrated that the processor possesses sufficient backend resources to sustain speculative execution in runahead mode. Therefore, we can simply use *nop* instructions, which do not consume backend resources, to validate the effectiveness of increasing the number of ROB entries through runahead execution. As shown in Fig.10, we measure the available ROB entries in three execution scenarios, including two cases with extended SEW:

- ① Without the runahead execution mechanism, the processor flushes the data “ x ” once.
- ② With the runahead execution mechanism, the processor flushes the data “ x ” once.
- ③ With the runahead execution mechanism, the processor flushes the data “ x ” multiple times. The condition that triggers runahead execution is set as a data cache miss, which theoretically allows an attacker to prolong the duration of the processor being in runahead mode by flushing the data multiple times. However, due to thread scheduling, the attacker must wait until all instructions in the ROB have retired before immediately flushing x and repeating




Figure 10. The code snippets for measuring the ROB size.




Figure 11. The access time of the probe array on a no-runahead machine and a runahead machine.

this process. And other processes cannot preempt the attacker in the meantime. Therefore, in real-world execution scenarios, case ③ is a probabilistic event. Our purpose is to point out that there is a possibility for further increasing the size of SEW.

Based on our measurements, the respective values for N_1 , N_2 , and N_3 are 255, 480, and 840, while ROB has 256 entries. Thus, it is evident that the number of executable instructions in runahead mode exceeds the capacity of ROB. To verify this advantage of the SPECRUN, we modify the code shown in Fig.8 by inserting multiple *nop* instructions in line 4 to ensure that the access instruction for secret data (line 5) appears outside the original ROB entries. We then execute the modified code on a no-runahead machine and a runahead machine, respectively. As shown in Fig.11, the experimental results exhibit that no significant latency drop is observed on the no-runahead machine, meaning the absence of data leakage. However, on the runahead machine, data leakage can still be observed at index 127. Hence, this proves that runahead-based speculative execution attacks can pose more significant harm, introducing the risk of data leakage to initially secure code.

6 A Secure Runahead Execution Mechanism

While the processor enters runahead mode, it speculatively executes subsequent instructions and loads data into the cache, creating a natural attack surface. We introduce a Speculative Load cache (SL cache) to address this vulnerability to store data loaded during runahead execution. Functioning as an “L0 cache”, the processor first accesses the SL cache upon exiting runahead mode. If a cache miss occurs, the data is searched from the L1 cache as usual. If the data exists in the SL cache, additional checks are required to determine whether it can be loaded into the L1 cache. To this end, we utilize the concept of Unsafe Speculative Load (USL) from InvisiSpec [18] to guide data placement. However, since branches in SPECRUN cannot be resolved, we restrict the definition of USL to the loads related to secret.

To be more specific, we employ taint tracking to identify the USLs during runahead execution. We introduce the tags B_{tag} and IS in the SL Cache. The variable B_n is defined as the n th branch, and $B_{tag} = B_{n,m}$ indicates that the load is the m th USL (i.e., taint-related load) within the scope of the n th branch. For the untainted load within the scope of B_n , B_{tag} is represented as $B_{n,0}$, while for the load outside the branch scope, B_{tag} is 0. In addition, the IS indicates whether the output of an instruction is associated with the taint of B_n . If the instruction’s input depends on tainted data within the

(a) machine code

```

rA = &A, rB = &B, rC = &C, rD =
&D, rE = &E, rF = &F, rG = &G,
rH = &H, rX = addr1, rY = addr2
if (rX < size_1){ //the start of B1
    ... //outer branch
    load r0 ← (rA)
    r1 = rB + rX //tainted data
    load r2 ← (r1) //tainted load
    r3 = rC * r2
    if (rY < size_2){ //the start of B2
        ... //inner branch
        r4 = rD - rY //tainted data
        load r5 ← (r4) //tainted load
        r6 = r5 + r2
        load r7 ← (r6) //tainted load
    } //the end of B2
    r8 = r3 - rE
    load r9 ← (r8) //tainted load
} //the end of B1
r10 = rF + r9
load r11 ← (r10) //tainted load
r12 = rG * r7
load r13 ← (r12) //tainted load
load r14 ← (rH)

```

(b) taint tracking

instr	output	input	B _{tag}	IS
branch		
...		
load	r0	rA	B _{1,0}	0
add	r1	rB, rX		B ₁
load	r2	r1	B _{1,1}	B ₁
mul	r3	rC, r2		B ₁
branch		
...		
sub	r4	rD, rY		B ₂
load	r5	r4	B _{2,1}	B ₂
add	r6	r5, r2		B ₁ , B ₂
load	r7	r6	B _{2,2}	B ₂
sub	r8	r3, rE		B ₁
load	r9	r8	B _{1,2}	B ₁
add	r10	rF, r9		B ₁
load	r11	r10	0	B ₁
mul	r12	rG, r7		B ₂
load	r13	r12	0	B ₂
load	r14	rH	0	0

Figure 12. The process of marking B_{tag} and IS . The processor encounters B_2 during the matching process for B_{1e} , indicating that branches B_1 and B_2 are nested. Load instructions are **red** and tainted registers are **blue**. scope of B_n , IS is set to B_n ; otherwise, IS is set to 0, meaning the load is safe. In this scheme, B_{tag} determines the security of the loads within the branch scope, while IS marks all potential USLs related to tainted data. However, identifying all tainted loads faces two challenges. First, nested branches need to be identified. Second, the taint-related loads outside the branch scope must be marked. The following is the specific solution to address these challenges.

The processor utilizes B_{ns} and B_{ne} to represent the start and end addresses of branch B_n based on the compiled code. If the processor encounters a new branch, such as B_y before matching B_{xe} , it means that B_y is an inner branch of B_x . Consequently, the matching order is changed to first match the end address of the inner branch, followed by that of the outer branch. When the processor needs to match B_{ne} , the loads during this period are marked as $B_{n,M}$ ($M \in [0, m]$). With these guidelines, we can effectively record the nested relationships of branches and mark the loads with $B_{n,m}$. Meanwhile, we determine the IS based on the tainted data propagation and the instruction inputs’ corresponding branch information. The process of marking B_{tag} and IS is shown in Fig.12.

Additionally, to mitigate the negative performance impact of unnecessary SL cache accesses, we add a counter C to record the number of data entries residing in the SL cache. Once C is reduced to 0, the processor ceases accessing the SL cache and instead loads data through the regular memory subsystem. The process of loading data by the processor after exiting the runahead mode is outlined in Algorithm 1. Concretely, we choose different methods to reduce the value of C depending on the data type:

1) If the data retrieved from the SL cache is considered safe ($IS = 0$), it is placed in the L1 cache, removed from the SL cache, and the value of C is decremented by 1.

2) In the case of USL ($IS \neq 0$), the processor waits for branch resolution. If the branch prediction is incorrect, the pre-stored nested branch information is utilized to determine which loads to remove. Specifically, the processor uses IS to delete the loads associated with the incorrectly executed branch itself and its inner branches. After that, the value of C is reduced based on the number of deleted entries. Otherwise, the execution is the same as above.

Algorithm 1: The load process after exiting runahead mode

Input: $S[]$, D_{SL} , d , B_n , B_{tag} , IS , C

```

1   $S[]$ : An array that stores the branches judged to be correctly
   predicted
2   $D_{SL}$ : The data fetched from the SL cache
3   $d$ : The number of entries in the SL cache that need to be deleted
   due to USL
4  processor encounters a load instruction;
5  if  $C \neq 0$  then
6      load data  $D_{SL}$  from the SL cache;
7      if  $B_{tag} \neq 0$  then
8          extract the  $B_n$  information;
9          if  $B_n$  is not in  $S[]$  then
10             wait the resolution of  $B_n$ ;
11             if the prediction of  $B_n$  is correct then
12                  $S[] \leftarrow B_n$ ;
13                  $D_{SL}$  is loaded to L1 cache;
14                  $C = C - 1$ ;
15             else
16                 use  $IS$  to delete entries related to  $B_n$ ;
17                  $C = C - d$ ;
18             else
19                  $D_{SL}$  is loaded to L1 cache;
20                  $C = C - 1$ ;
21         else
22              $D_{SL}$  is loaded to L1 cache;
23              $C = C - 1$ ;
24     else
25         loads data from the regular cache;
26         switch execution paths so that subsequent loads before the next
         round of runahead execution no longer access the SL cache;

```

In addition to designing a secure runahead execution scheme, we can also consider imposing restrictions on existing mechanisms. For example, we can nullify the impact of branches on instruction execution within the runahead interval. Once a branch predicate is identified as associated with a stalling load, the branch is skipped rather than unresolved. In summary, before deploying runahead execution, it is crucial to balance security and performance carefully. While effective, the countermeasures may lead to increased overhead and degraded performance to some extent.

7 Related Work

A large number of attack variants have been proposed based on optimization mechanisms of commercial processors such as out-of-order execution, branch prediction, store-to-load forwarding, memory disambiguation, speculative dependency checking, prefetching, fault handling, microcode assist, and others. Microarchitectural vulnerabilities have sparked widespread concern, as they leak arbitrary confidential information in a reliable and fine-grained manner. What is even more concerning is that as Moore's Law slows down, future processors may adopt various technologies to enhance performance, which may also introduce unknown security issues. To address this impending security crisis, a trend of the current studies is to perform prospective security analysis on optimization mechanisms that may be deployed in the future. Several recent works have focused on this topic.

[6] proposes a conceptual framework to study seven classes of microarchitectural optimization with security implications, ranging from value locality to compressibility to prediction to prefetching, then provides a PoC implementation of silent stores and data

memory-dependent. [7] offers the first security analysis of cache compression and introduces a new concealing side channel, as simply storing data transmits information about it. [8] proposes a new security attack on value predictor. Furthermore, the authors developed a model for analyzing value predictor attacks, based on which they found 12 attack variants.

Our proposed SPECRUN leverages the potential of runahead execution, an optimization mechanism aimed at mitigating the performance impact of long-latency accesses. SPECRUN has simple trigger conditions and high applicability compared to previous works. What sets it apart is the utilization of a mixture of optimization mechanisms, i.e., nested branch prediction within runahead execution. This allows SPECRUN to eliminate the limitation on the number of transient instructions imposed by the ROB size, leading to more severe security threats.

8 Conclusion

Runahead execution is a promising microarchitectural technique designed to mitigate the memory wall. However, its security has yet to be evaluated prior to this work. As demonstrated in this study, runahead execution is also threatened by transient execution attacks. We propose a novel attack named SPECRUN, which takes advantage of the microarchitectural design of unresolved branch prediction in runahead processors and eliminates the limitation on the number of transient instructions. In light of this, we also design a secure runahead execution scheme based on invisible and tainted tracking. We hope that this research will prompt careful consideration of security before implementing optimization mechanisms in future commercial processors.

Acknowledgment

Shen and Zhang are supported by the National Natural Science Foundation of China under Grants No. 62122023 and U20A20202, the Postgraduate Scientific Research Innovation Project of Hunan Province under Grant No. CX20210406.

References

- [1] P. Kocher et al., "Spectre Attacks: Exploiting Speculative Execution," in *S&P*, pp. 1-19, 2019.
- [2] M. Lipp et al., "Meltdown: Reading Kernel Memory from User Space," in *USENIX Security*, pp. 973-990, 2018.
- [3] Y. Chen et al., "AfterImage: Leaking Control Flow Data and Tracking Load Operations via the Hardware Prefetcher," in *ASPLOS*, pp. 16-32, 2023.
- [4] S. Islam et al., "SPOILER: Speculative Load Hazards Boost Rowhammer and Cache Attacks," in *USENIX Security*, pp. 621-637, 2019.
- [5] D. Skarlatos et al., "MicroScope: enabling microarchitectural replay attacks," in *ISCA*, pp. 318-331, 2019.
- [6] J. R. S. Vicarte et al., "Opening Pandora's Box: A Systematic Study of New Ways Microarchitecture Can Leak Private Data," in *ISCA*, pp. 88-105, 2021.
- [7] Po-An Tsai et al., "Safecracker: Leaking Secrets through Compressed Caches," in *ASPLOS*, pp. 1125-1140, 2020.
- [8] S. Deng et al., "New Predictor-Based Attacks in Processors," in *DAC*, pp. 697-702, 2021.
- [9] O. Mutlu et al., "Runahead execution: an alternative to very large instruction windows for out-of-order processors," in *HPCA*, pp. 129-140, 2003.
- [10] A. Naithani et al., "Vector Runahead," in *ISCA*, pp. 195-208, 2021.
- [11] M. Hashemi et al., "Filtered Runahead Execution with a Runahead Buffer," in *MIRCO*, pp. 358-369, 2015.
- [12] M. Hashemi et al., "Continuous Runahead: Transparent Hardware Acceleration for Memory Intensive Workloads," in *MIRCO*, pp. 1-12, 2016.
- [13] A. Naithani et al., "Precise Runahead Execution," in *HPCA*, pp. 397-410, 2020.
- [14] G. Maisuradze et al., "ret2spec: Speculative Execution Using Return Stack Buffers," in *CCS*, pp. 2109-2122, 2018.
- [15] M. H. I. Chowdhury et al., "BranchSpec: Information Leakage Attacks Exploiting Speculative Branch Instruction Executions," in *ICCD*, pp. 529-536, 2020.
- [16] Intel, "Intel® 64 and IA-32 Architectures Software Developer's Manual," <https://www.intel.com/content/www/us/en/developer/articles/technical/intel-sdm.html>, 2019.
- [17] R. Ubal et al., "Multi2Sim: a simulation framework for CPU-GPU computing," in *PACT*, pp. 335-344, 2021.
- [18] M. Yan et al., "InvisiSpec: Making Speculative Execution Invisible in the Cache Hierarchy," in *MIRCO*, pp. 428-441, 2018.