

Title: F1: A Fast and Programmable Accelerator for Fully Homomorphic Encryption

(2021 MICRO)

Introduction

Challenges

- Complex operations on long vectors: modular arithmetic, several thousand elements
- Regular computation: all operations are known ahead of time, VLIM
- Challenging data movement: large amounts (tens of MBs) of data; encrypting data increase its size(50X); data in long vectors

Introduction

Contributions

- F1 features an explicitly managed on-chip memory hierarchy, with a heavily bank scratchpad and distributed reg files
- **F1 uses mechanisms to decoupled data movement and hide access latencies by loading data far ahead of its use**
- F1 uses new scheduling algorithms that maximize reuse and make the best out of limited memory bandwidth
- F1 used few functional units with high throughput that reduces the amount of data

Background

FHE programming model and operations

- element-wise
- addition (mod t)
- multiplication (mod t)
- a small set of particular vector permutations.

BGV implementation overview

Data types:

$$a = a_0 + a_1x + \dots + a_{N-1}x^{N-1} \in R_t$$

Each plaintext is encrypted into a ciphertext consisting of two polynomials of N integer coefficients modulo some $Q \gg t$. Each ciphertext polynomial is a member of R_Q .

Encryption and decryption:

secret key: $s \in R_Q$. To encrypt a plaintext $m \in R_t$, one samples a uniformly random $a \in R_Q$, an error (or noise) $e \in R_Q$ with small entries, and computes the ciphertext ct as

$$ct = (a, b = as + te + m)$$

Ciphertext $ct = (a, b)$ is decrypted by recovering $e' = te + m = b - as \pmod{Q}$, and then recovering $m = e' \pmod{t}$. Decryption is correct as long as e' does not “wrap around” modulo Q , i.e., its coefficients have magnitude less than $Q/2$.

Homomorphic operations

- addition

$$ct_0 = (a_0, b_0) \text{ and } ct_1 = (a_1, b_1)$$

$$ct_{add} = ct_0 + ct_1 = (a_0 + a_1, b_0 + b_1)$$

- multiplication

$$ct_{\times} = (l_2, l_1, l_0) = (a_0a_1, a_0b_1 + a_1b_0, b_0b_1)$$

$$(u_1, u_0) = \textit{KeySwitch}(I_2)$$

$$ct_{mul} = (I_1 + u_1, I_0 + u_0)$$

- permutations

There are N automorphisms, denoted $\sigma_k(a)$ and $\sigma_{-k}(a)$ for all positive odd $k < N$. Specifically, $\sigma_k(a) : a_i \mapsto (-1)^s a_{ik} \pmod N$ for $i = 0, \dots, N-1$

where $s = 0$ if $ik \pmod{2N} < N$, and $s = 1$ otherwise.

1. compute an automorphism on the ciphertext polynomials:

$$ct_\sigma = (\sigma_k(a), \sigma_k(b))$$

$$2. ct_{perm} = (u_1, \sigma_k(b) + u_0) \text{ where } (u_1, u_0) = \text{KeySwitch}(\sigma_k(a))$$

Noise growth and management

Different operations induce different noise growth: addition and permutations cause little growth, but multiplication incurs much more significant growth. So, to a first order, the amount of noise is determined by **multiplicative depth**, i.e., the longest chain of homomorphic multiplications in the computation.

Noise forces the use of a large ciphertext modulus Q . For example, an FHE program with multiplicative depth of 16 needs Q to be about 512 bits. The noise budget, and thus the tolerable multiplicative depth, grow linearly $\log Q$

- Bootstrapping:
strength: enable FHE computations of unbounded depth; remove noise from a ciphertext without access to the secret key
weakness: need a large noise budget (large Q)
- Modulus switching:
rescales ciphertexts from modulus Q to a modulus Q' .
To execute an multiplicative depth 16, we start with a 512 bit modulus Q . Before multiplicatino, switch to a modulus that is 32 bits shorter.

Security and parameters

dimension N and modulus Q

$N/\log Q$ must be above a certain level for sufficient security.

Algorithmic insights and optimizations

- Fast polynomial multiplication via NTTs
- Avoiding wide arithmetic via Residue Number System(RNS)

Architectural analysis of FHE

Three input: a polynomial x (store in L residue polynomials), two key-switch hint matrices $ksh0, ksh1$.

Inputs and outputs are in the NTT domain; only $y[i]$ are in coefficient form.

```
1  def keySwitch(x: RVec[L],
2              ksh0: RVec[L][L], ksh1: RVec[L][L]):
3      y = [INTT(x[i], qi) for i in range(L)]
4      u0: RVec[L] = [0, ...]
5      u1: RVec[L] = [0, ...]
6      for i in range(L):
7          for j in range(L):
8              xqj = (i == j) ? x[i] : NTT(y[i], qj)
9              u0[j] += xqj * ksh0[i, j] mod qj
10             u1[j] += xqj * ksh1[i, j] mod qj
11     return (u0, u1)
```

Listing 1: Key-switch implementation. **RVec** is an N -element vector of 32-bit values, storing a single RNS polynomial in either the coefficient or the NTT domain.

Computation vs. data movement

- L^2 NTTs, $2L^2$ multiplications, $2L^2$ additions of N -element vectors
- In RNS, the rest of a homomorphic multiplication is $4L$ multiplications and $3L$ additions

$$L = 16, N = 16K$$

each RNS polynomial is 64KB, each polynomial is 1MB, each ciphertext is 2MB, key switch hints is 32MB.

key switching demand 10TB/s of memory and bandwidth.

- Performance requirement:
 - (1) decouples data movement from computation, as demand misses during frequent key-switches would tank performance
 - (2) implements a large amount of on-chip storage (over 32 MB in our example) to allow reuse across entire homomorphic operations
- Functionality requirements:

Programmable FHE accelerators must support a wide range of parameters, both N (polynomial/vec-tor sizes) and L (number of RNS polynomials, i.e., number of 32-bit prime factors of Q). While N is generally fixed for a single program, L changes as modulus switching sheds off polynomials.

F1 ARCHITECTURE

Vector processing with specialized functional units

FUs process vectors of configurable length N using a fixed number of vector lanes E .

- 128 lanes
- N from 1024 to 16384
- pipelined, throughput: $E = 128$ elements/cycle




Figure 2: Overview of the F1 architecture.

Compute clusters:

- 1 NTT, 1 automorphism
- 2 multipliers
- 2 adders
- a banked register file

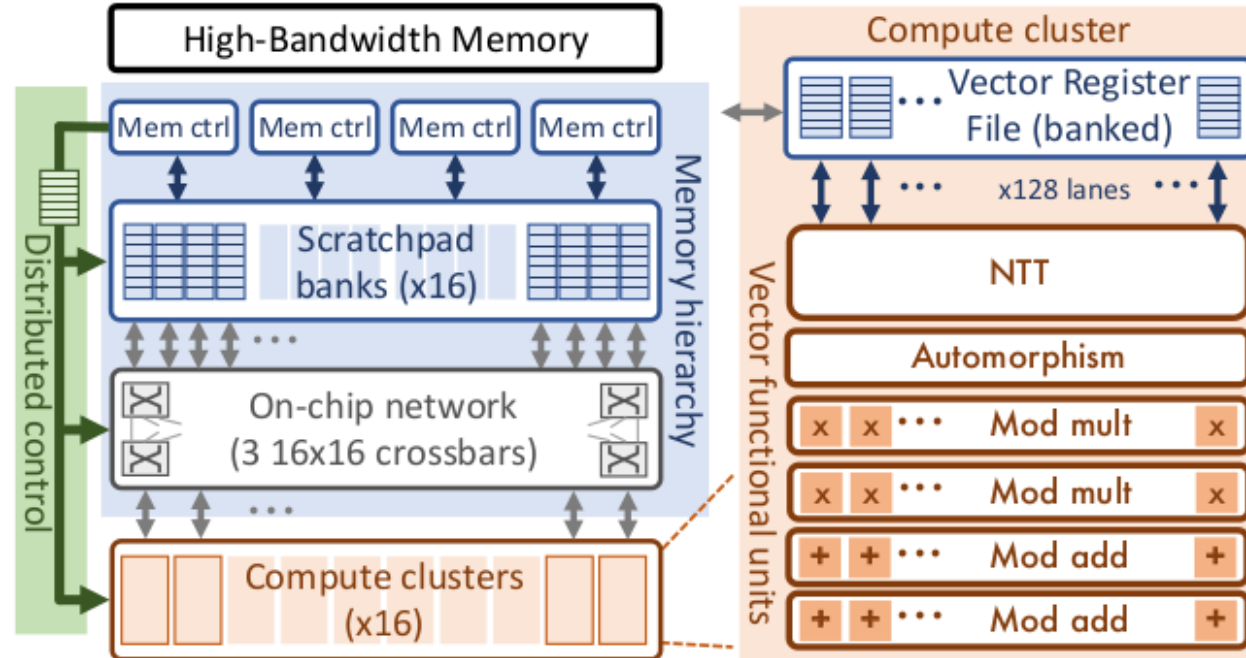


Figure 2: Overview of the F1 architecture.

Memory system:

- a large, heavily banked scratchpad (64 MB across 16 banks)
- scratchpad interfaces with both high-bandwidth off-chip memory (HBM2) and with compute clusters through an on-chip network




Figure 2: Overview of the F1 architecture.

Static scheduling(programs are regular):

- VLIW processors?
- FUs: no stalling logic
- Memory: no conflicts
- On-chip network: use switch change configuration




Figure 2: Overview of the F1 architecture.

Distribute control:

- independent instruction stream: programs have loops, unroll them avoid branches, and compile programs into linear sequences of instructions




Figure 2: Overview of the F1 architecture.

F1 ARCHITECTURE

Register file design:

use an 8-banked element-partitioned register file design that leverages long vectors: each vector is striped across banks, and each FU cycles through all banks over time, using a single bank each cycle




Figure 2: Overview of the F1 architecture.

SCHEDULING DATA AND COMPUTATION

- Compiler: orders high level operations to maximize reuse and translates the program into a DFG

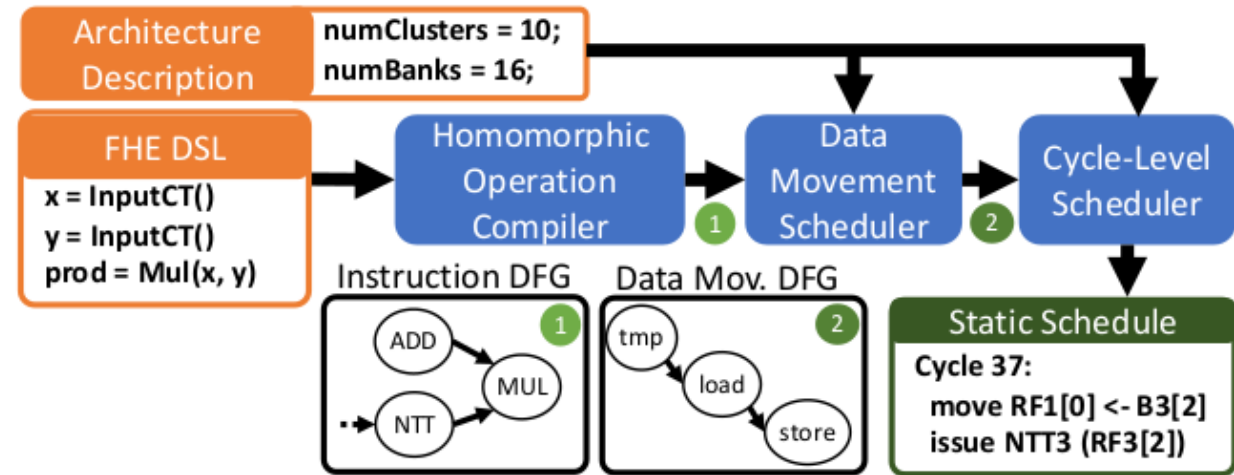


Figure 3: Overview of the F1 compiler.

- DM Scheduler: transfer between main memory and the scratchpad to achieve decoupling and maximize reuse

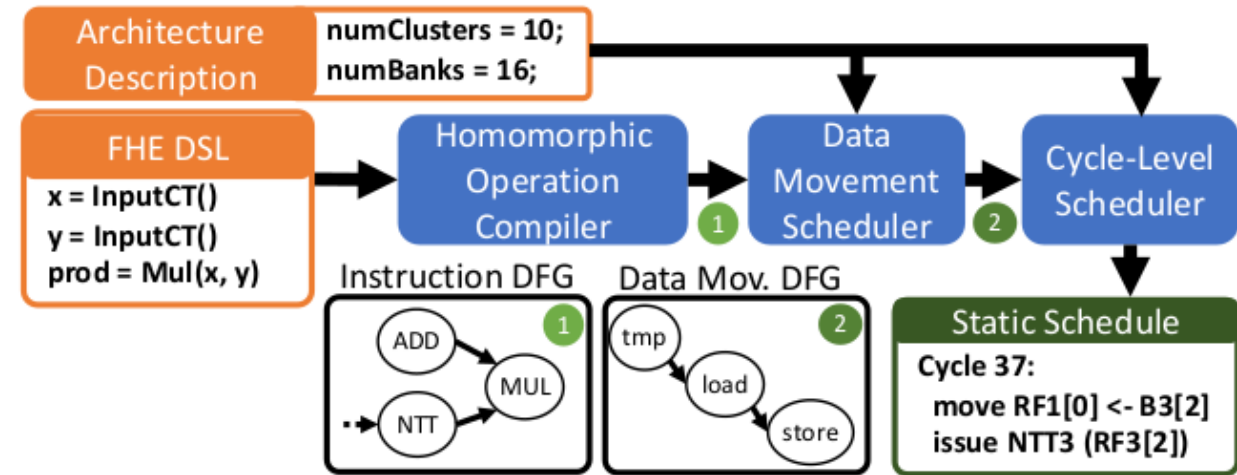


Figure 3: Overview of the F1 compiler.

- CL Scheduler: determine the exact cycles of all operations and produces the instruction strams for all components

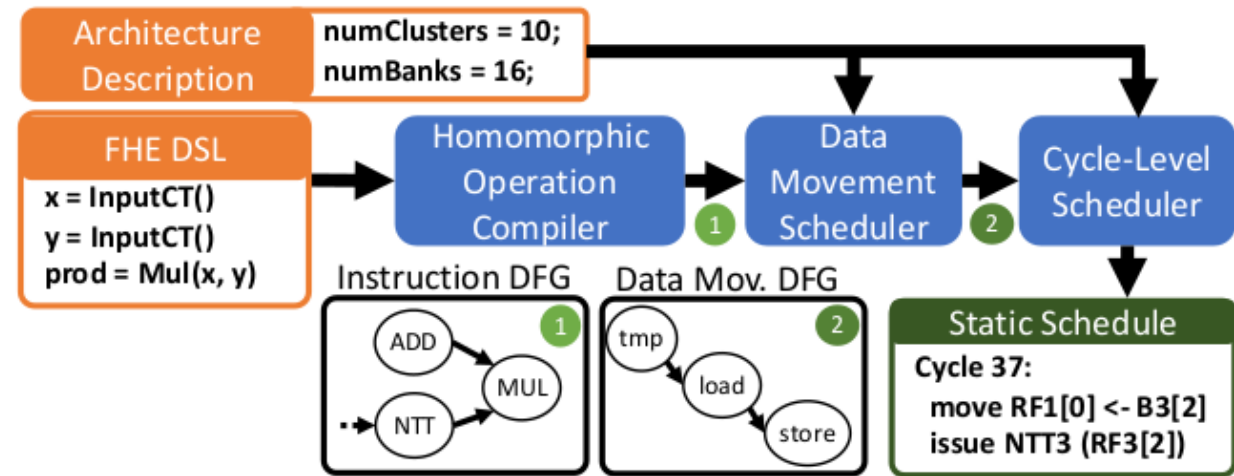


Figure 3: Overview of the F1 compiler.

Translating the program to a dataflow graph

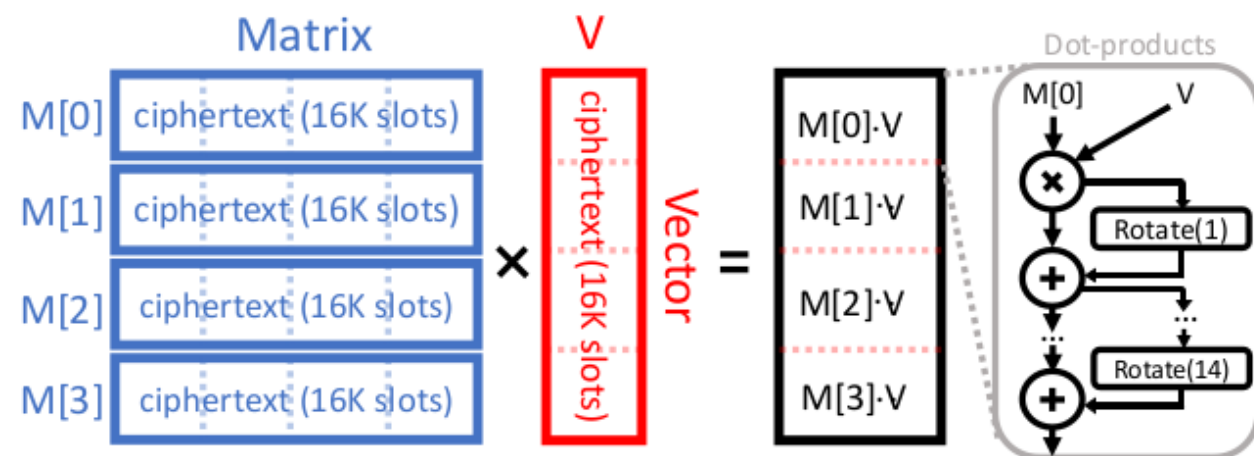


Figure 4: Example matrix-vector multiply using FHE.

```
1  p = Program(N = 16384)
2  M_rows = [ p.Input(L = 16) for i in range(4) ]
3  output = [ None for i in range(4) ]
4  V = p.Input(L = 16)
5
6  def innerSum(X):
7      for i in range(log2(p.N)):
8          X = Add(X, Rotate(X, 1 << i))
9      return X
10
11 for i in range(4):
12     prod = Mul(M_rows[i], V)
13     output[i] = innerSum(prod)
```

Listing 2: $(4 \times 16K)$ matrix-vector multiply in F1's DSL.

Compiling homomorphic operations

It clusters operations to improve reuse and translates them down to instruction.

- Ordering: maximize the reuse of key switch hints(line 8)(line 12)
- Translation: minimize the amount of instructions intermediates

Scheduling data transfers

- data transfers decoupled from computation
- minimize off-chip data transfers
- achieve good parallelism

It does not consider on-chip data movement, and simply treats all functional units as being directly connected to the scratchpad.

It considers instructions ready if their inputs are available in the scratchpad, and follows instruction priority among ready ones. To schedule loads, we assign each load a priority

$$p(load) = \max\{p(u) | u \in users(load)\}$$

then greedily issue loads as bandwidth becomes available. When issuing an instruction, we must ensure that there is space to store its result. We can often replace a dead value.

Cycle-level scheduling(constrained by its input schedule's off-chip data movement)

- distribute computation across clusters and manage reg file and on-chip transfer
- add loads or stores in this stage
- move loads to their earliest possible issue cycle to avoid stalls on missing operands

FUNCTIONAL UNITS

Automorphism unit

how automorphism σ_3 is applied to a residue polynomial with $N = 16$ and $E = 4$ elements/cycle.




Figure 5: Applying σ_3 on an RNS polynomial of four 4-element chunks by using only permutations local to chunks.

Automorphism unit

Given a residue polynomial of $N = G \cdot E$ elements, the automorphism unit first transpose applies the column permutation to each E -element input. Then, it feeds this to a transpose unit that reads in the whole residue polynomial interpreting it as a $G \times E$ matrix, and produces its transpose $E \times G$.




Figure 6: Automorphism unit.

Transpose unit

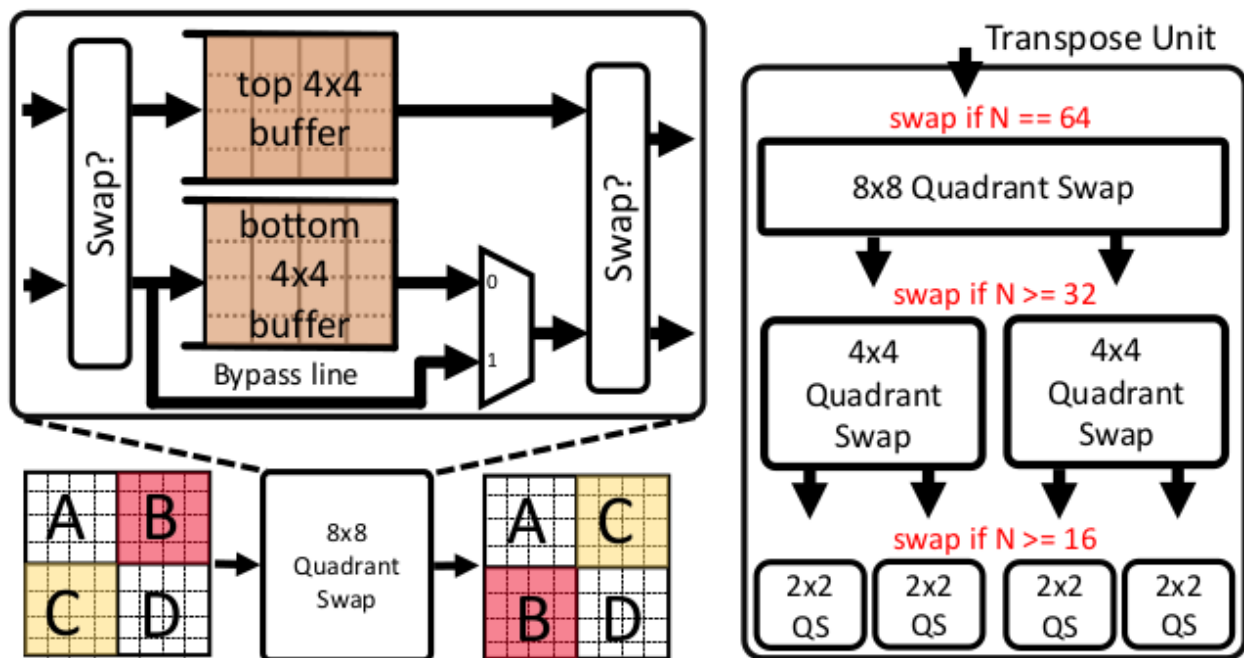


Figure 7: Transpose unit (right) and its component quadrant-swap unit (left).

Transpose unit: Our *quadrant-swap transpose* unit transposes an $E \times E$ (e.g., 128×128) matrix by recursively decomposing it into quadrants and exploiting the identity

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix}^T = \begin{bmatrix} A^T & C^T \\ B^T & D^T \end{bmatrix}.$$

Four-step NTT unit




Figure 8: Example of a four-step NTT datapath that uses 4-point NTTs to implement 16-point NTTs.

Optimized modular multiplier

Multiplier	Area [μm^2]	Power [mW]	Delay [ps]
Barrett	5,271	18.40	1,317
Montgomery	2,916	9.29	1,040
NTT-friendly	2,165	5.36	1,000
FHE-friendly (ours)	1,817	4.10	1,000

Table 1: Area, power, and delay of modular multipliers.

F1 IMPLEMENTATION

Component	Area [mm ²]	TDP [W]
NTT FU	2.27	4.80
Automorphism FU	0.58	0.99
Multiply FU	0.25	0.60
Add FU	0.03	0.05
Vector RegFile (512 KB)	0.56	1.67
Compute cluster (NTT, Aut, 2× Mul, 2× Add, RF)	3.97	8.75
Total compute (16 clusters)	63.52	140.0
Scratchpad (16×4 MB banks)	48.09	20.35
3×NoC (16×16 512 B bit-sliced [58])	10.02	19.65
Memory interface (2×HBM2 PHYs)	29.80	0.45
Total memory system	87.91	40.45
Total F1	151.4	180.4

Table 2: Area and Thermal Design Power (TDP) of F1, and breakdown by component.

EXPERIMENTAL METHODOLOGY

- Modeled system:
a cycle-accurate simulator to execute F1 programs
activity-level energies from RTL synthesis to produce energy breakdowns
- Benchmarks:
Logistic regression: uses the HELR algorithm: 256 features, 256 samples, depth $L = 16$
Neural network: LoLa-MNIST, LoLa-CIFAR
DB Lookup: A BGV-encrypted query string is used to traverse an encrypted key-value store and return the corresponding value.

- Bootstrapping:
BGV: Sheriff and Peikert's algorithm
CKKS: non-packed CKKS bootstrapping
- Baseline systems:
F1 with a CPU system running the baseline programs (a 4-core, 8-thread, 3.5 GHz Xeon E3-1240v5)

EVALUATION

Performance

Benchmarks

Execution time (ms) on	CPU	F1	Speedup
LoLa-CIFAR Unencryp. Wghts.	1.2×10^6	241	5,011×
LoLa-MNIST Unencryp. Wghts.	2,960	0.17	17,412×
LoLa-MNIST Encryp. Wghts.	5,431	0.36	15,086×
Logistic Regression	8,300	1.15	7,217×
DB Lookup	29,300	4.36	6,722×
BGV Bootstrapping	4,390	2.40	1,830×
CKKS Bootstrapping	1,554	1.30	1,195×
gmean speedup			5,432×

*LoLa's release did not include MNIST with encrypted weights, so we reimplemented it in HELib.

Table 3: Performance of F1 and CPU on full FHE benchmarks: execution times in milliseconds and F1's speedup.

Microbenchmarks

	$N = 2^{12}, \log Q = 109$			$N = 2^{13}, \log Q = 218$			$N = 2^{14}, \log Q = 438$		
	F1	vs. CPU	vs. HEAX _{σ}	F1	vs. CPU	vs. HEAX _{σ}	F1	vs. CPU	vs. HEAX _{σ}
NTT	12.8	17,148×	1,600×	44.8	10,736×	1,733×	179.2	8,838×	1,866×
Automorphism	12.8	7,364×	440×	44.8	8,250×	426×	179.2	16,957×	430×
Homomorphic multiply	60.0	48,640×	172×	300	27,069×	148×	2,000	14,396×	190×
Homomorphic permutation	40.0	17,488×	256×	224	10,814×	198×	1,680	6,421×	227×

Table 4: Performance on microbenchmarks: F1’s reciprocal throughput, in nanoseconds per ciphertext operation (lower is better) and speedups over CPU and HEAX _{σ} (HEAX augmented with scalar automorphism units) (higher is better).

Architectural analysis

Data movement, Power consumption, Utilization over time


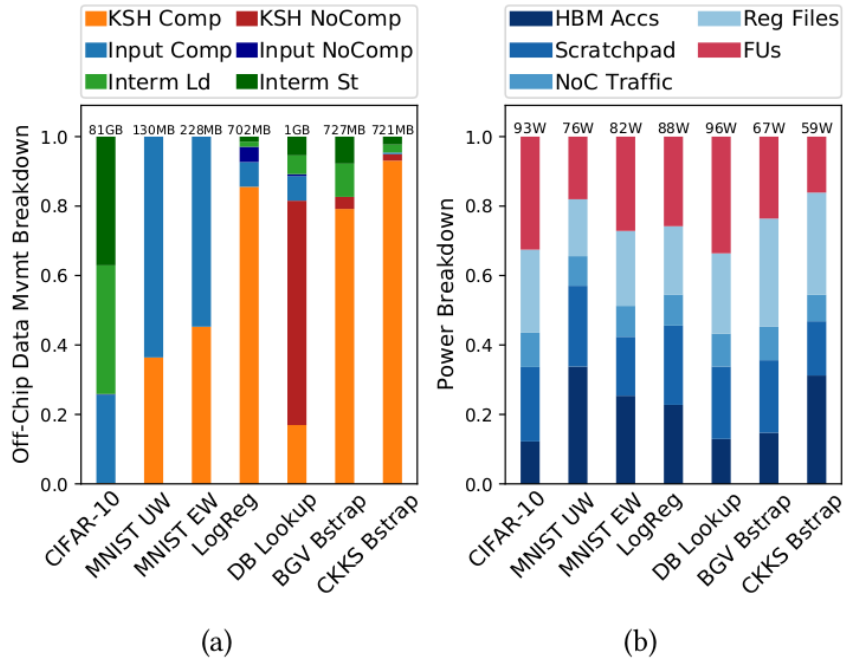


Figure 10: Functional unit and HBM utilization over time for the LoLa-MNIST PTW benchmark.

Sensitivity studies

Benchmark	LT NTT	LT Aut	CSR
LoLa-CIFAR Unencryp. Wghts.	3.5×	12.1×	—*
LoLa-MNIST Unencryp. Wghts.	5.0×	4.2×	1.1×
LoLa-MNIST Encryp. Wghts.	5.1×	11.9×	7.5×
Logistic Regression	1.7×	2.3×	11.7×
DB Lookup	2.8×	2.2×	—*
BGV Bootstrapping	1.5×	1.3×	5.0×
CKKS Bootstrapping	1.1×	1.2×	2.7×
gmean speedup	2.5×	3.6×	4.2×

*CSR is intractable for this benchmark.

Table 5: Speedups of F1 over alternate configurations: LT NTT/Aut = Low-throughput NTT/Automorphism FUs; CSR = Code Scheduling to minimize Register Usage [37].

Scalability




Figure 11: Performance vs. area across F1 configurations.